



Human-Centered Data & AI

Vinicius Caridá, Ph.D.



Data Science Manager, Itaú Unibanco

MBA Professor, FIAP

GDE – Machine Learning



@vinicius caridá



@vfcarida



@vinicius caridá



@vfcarida



@vinicius caridá



@vfcarida

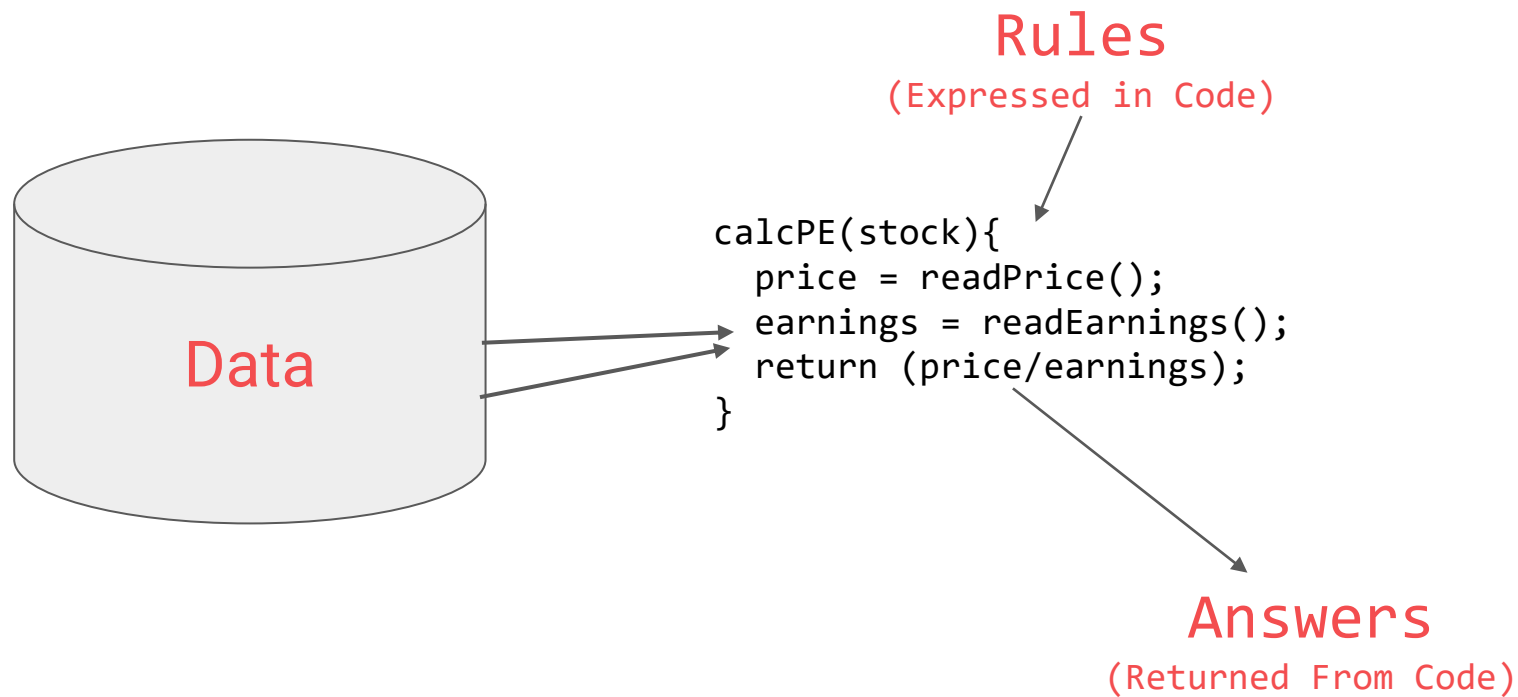



# Zero to Hero with TensorFlow

Vinicius Caridá

@vfcarida







```
if (ball.collide(brick)){  
    removeBrick();  
    ball.dx=-1*(ball.dx);  
    ball.dy=-1*(ball.dy);  
}
```







# Activity Recognition



```
if(speed<4){  
    status=WALKING;  
}
```





# Activity Recognition



```
if(speed<4){  
    status=WALKING;  
}
```



```
if(speed<4){  
    status=WALKING;  
} else {  
    status=RUNNING;  
}
```



# Activity Recognition



```
if(speed<4){  
    status=WALKING;  
}
```



```
if(speed<4){  
    status=WALKING;  
} else {  
    status=RUNNING;  
}
```



```
if(speed<4){  
    status=WALKING;  
} else if(speed<12){  
    status=RUNNING;  
} else {  
    status=BIKING;  
}
```



# Activity Recognition



```
if(speed<4){  
    status=WALKING;  
}
```



```
if(speed<4){  
    status=WALKING;  
} else {  
    status=RUNNING;  
}
```



```
if(speed<4){  
    status=WALKING;  
} else if(speed<12){  
    status=RUNNING;  
} else {  
    status=BIKING;  
}
```



```
// ????
```





# Activity Recognition



0101001010100101010  
1001010101001011101  
0100101010010101001  
0101001010100101010

Label = WALKING



1010100101001010101  
0101010010010010001  
0010011111010101111  
1010100100111101011

Label = RUNNING



1001010011111010101  
1101010111010101110  
1010101111010101011  
1111110001111010101

Label = BIKING



1111111111010011101  
0011111010111110101  
0101110101010101110  
1010101010100111110

Label = GOLFING  
(Sort of)



# Activity Recognition



0101001010100101010  
1001010101001011101  
0100101010010101001  
0101001010100101010

Label = WALKING



1010100101001010101  
0101010010010010001  
0010011111010101111  
1010100100111101011

Label = RUNNING



1001010011111010101  
1101010111010101110  
1010101111010101011  
1111110001111010101

Label = BIKING



1111111111010011101  
0011111010111110101  
0101110101010101110  
1010101010100111110

Label = GOLFING  
(Sort of)



# Activity Recognition



0101001010100101010  
1001010101001011101  
0100101010010101001  
0101001010100101010

Label = WALKING



1010100101001010101  
0101010010010010001  
0010011111010101111  
1010100100111101011

Label = RUNNING



1001010011111010101  
1101010111010101110  
1010101111010101011  
1111110001111010101

Label = BIKING



1111111111010011101  
0011111010111110101  
0101110101010101110  
1010101010100111110

Label = GOLFING  
(Sort of)




# The Machine Learning Paradigm

Make a Guess!



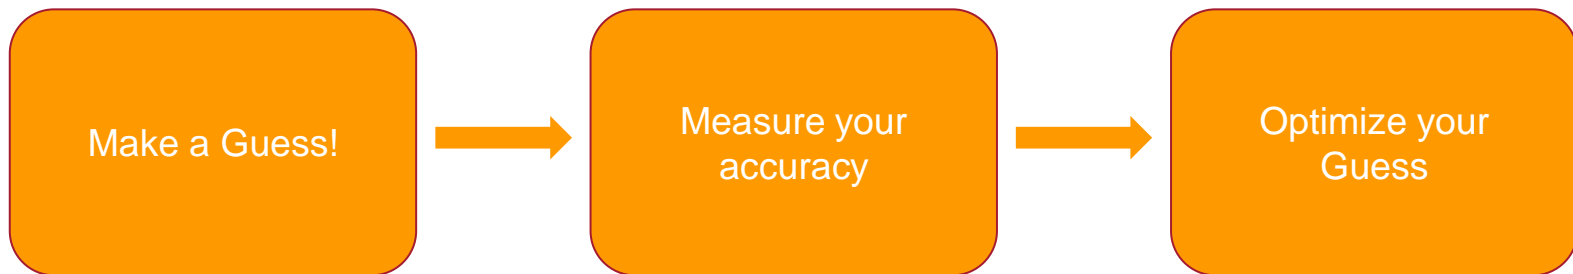


# The Machine Learning Paradigm





# The Machine Learning Paradigm






# The Machine Learning Paradigm



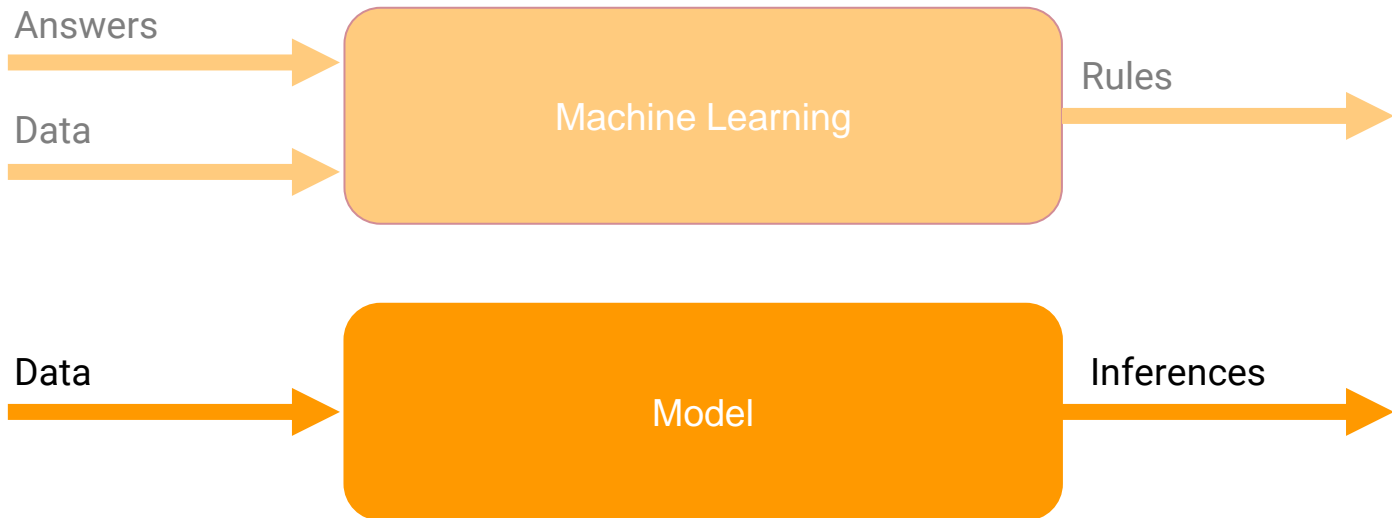


# The Machine Learning Paradigm





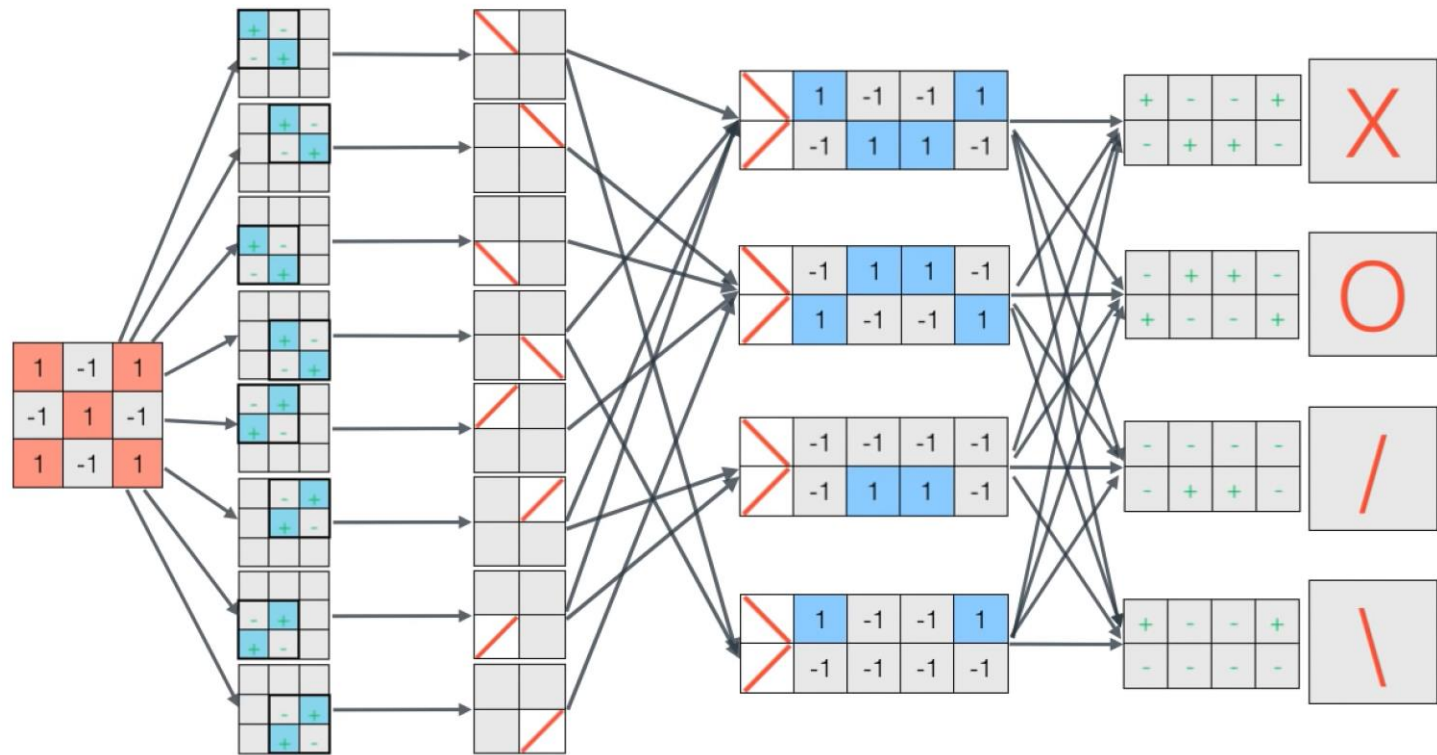
# The Machine Learning Paradigm





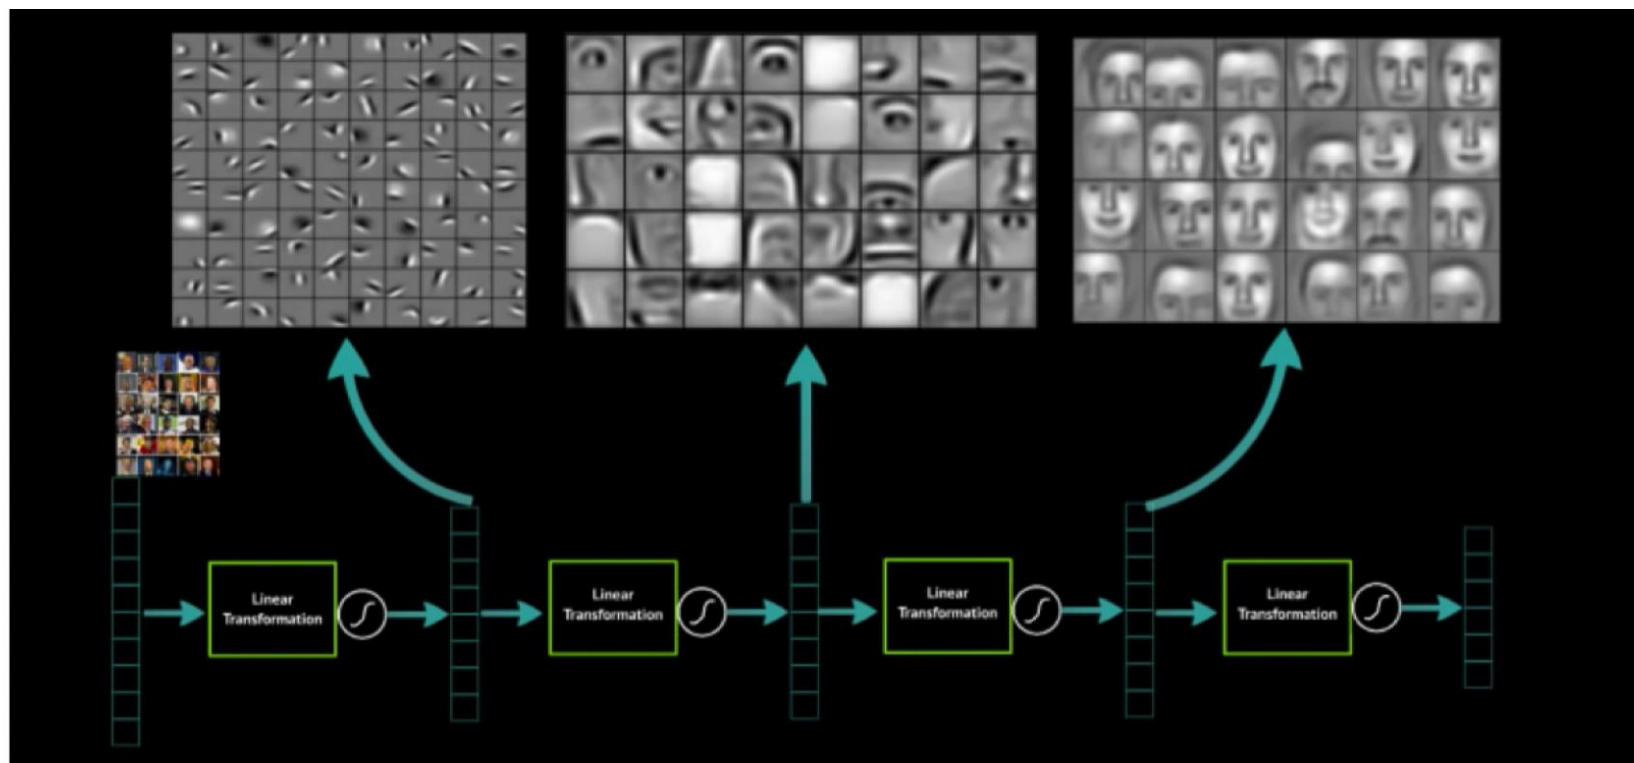
# Computer Vision

Using Machine Learning to understand images



Convolution Layer Pooling Layer

Fully Connected Layer























+






+




+








+




+




=


HUMAN



+




+




=


HUMAN



+




+

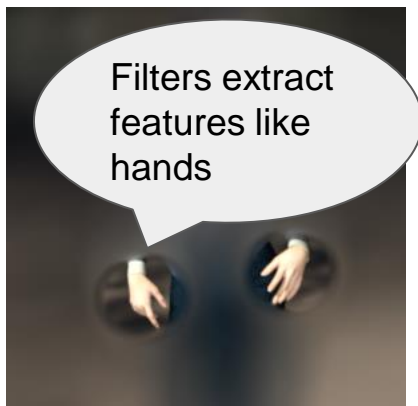


=


HORSE



+



+




=


HUMAN



+




+



=


HORSE



+




+



=

HUMAN




+




=

HORSE






Filters can then be combined with labels to make a prediction of the image contents...



Filters can then be combined with labels to make a prediction of the image contents...

The filters that match the label are learned over time!





Data



Model

Inferences





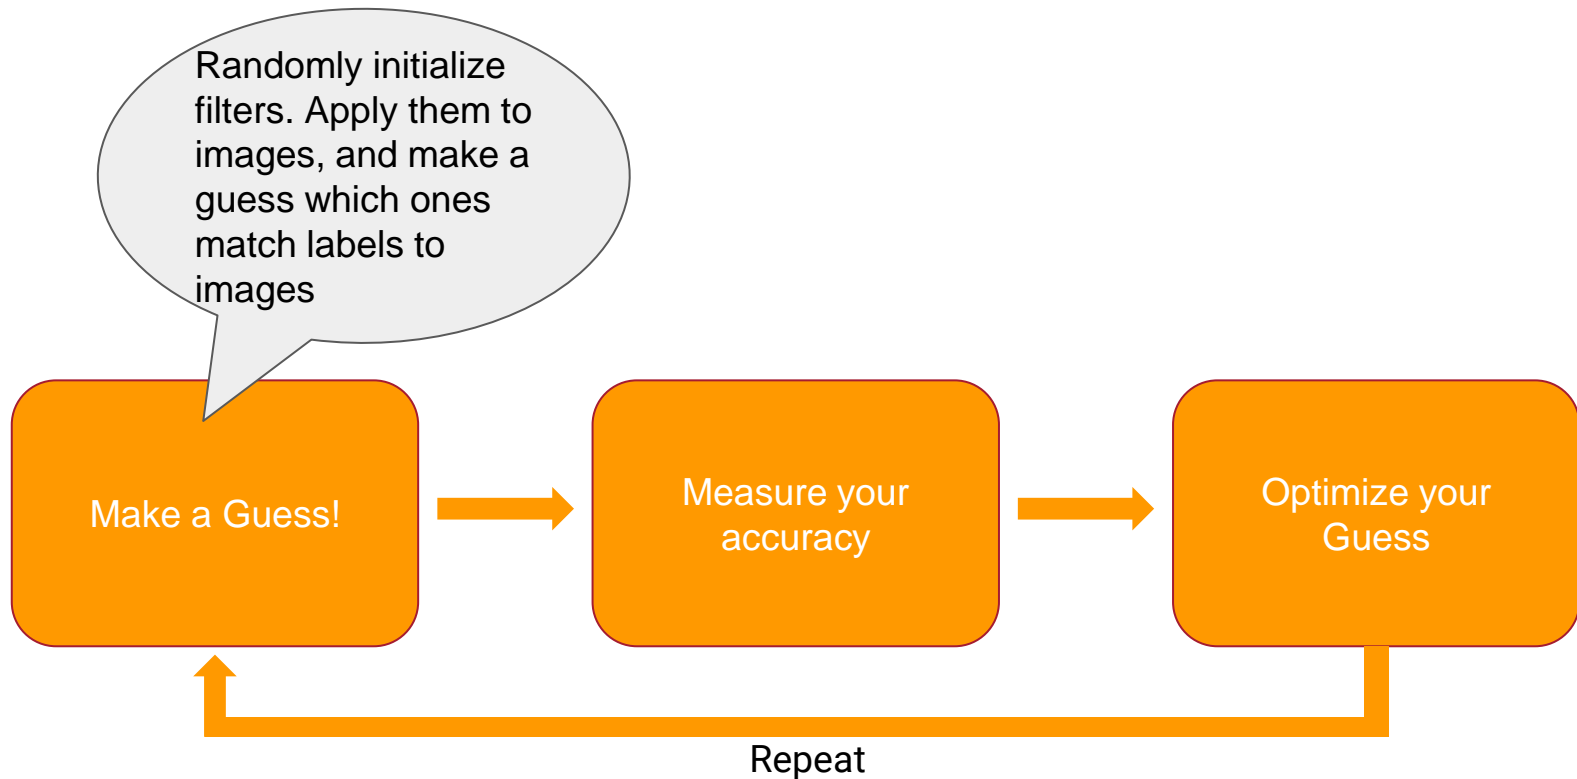


# The Machine Learning Paradigm



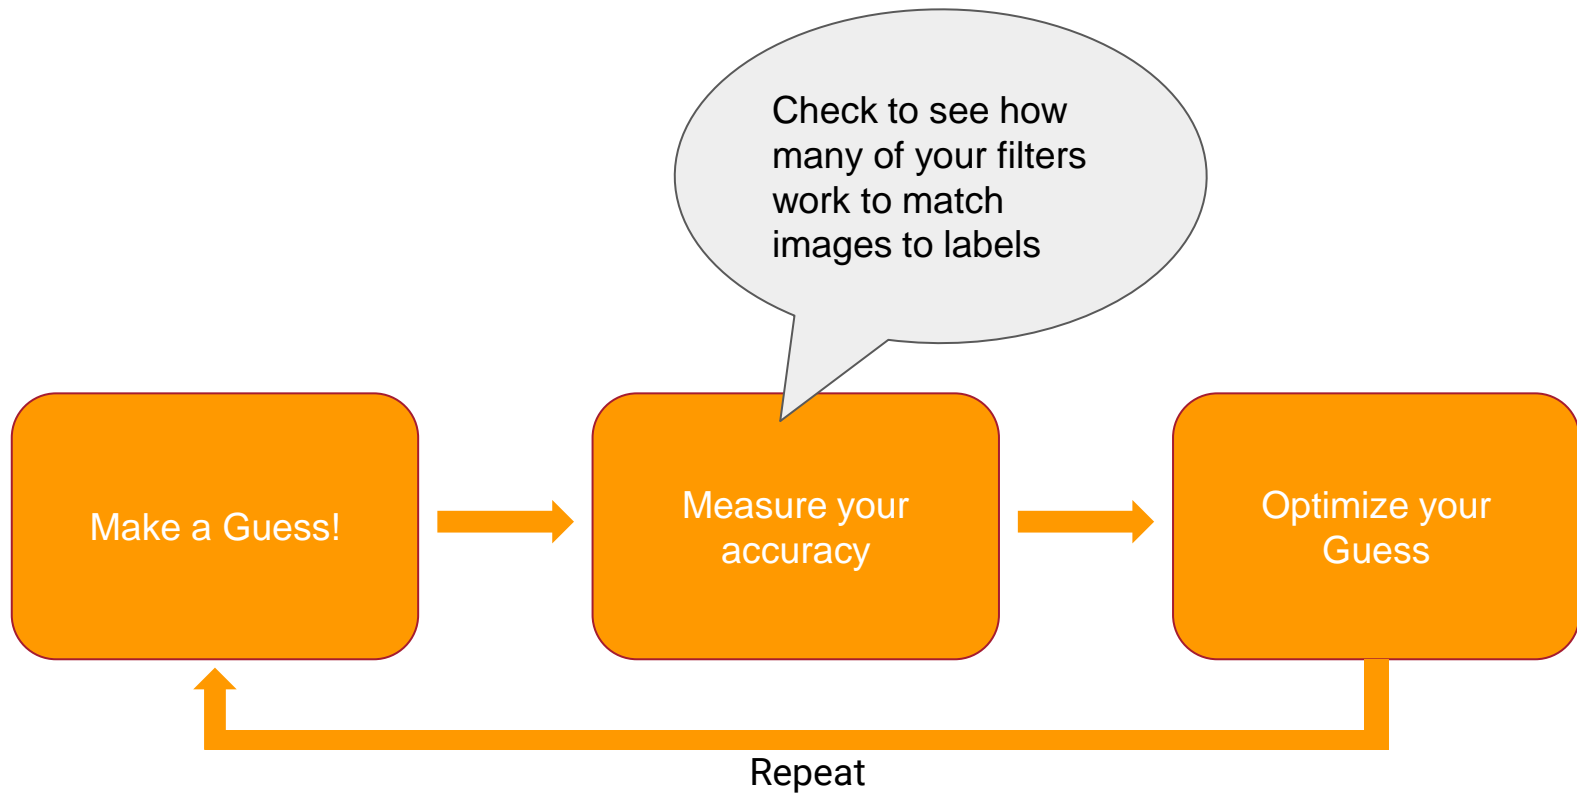


# The Machine Learning Paradigm



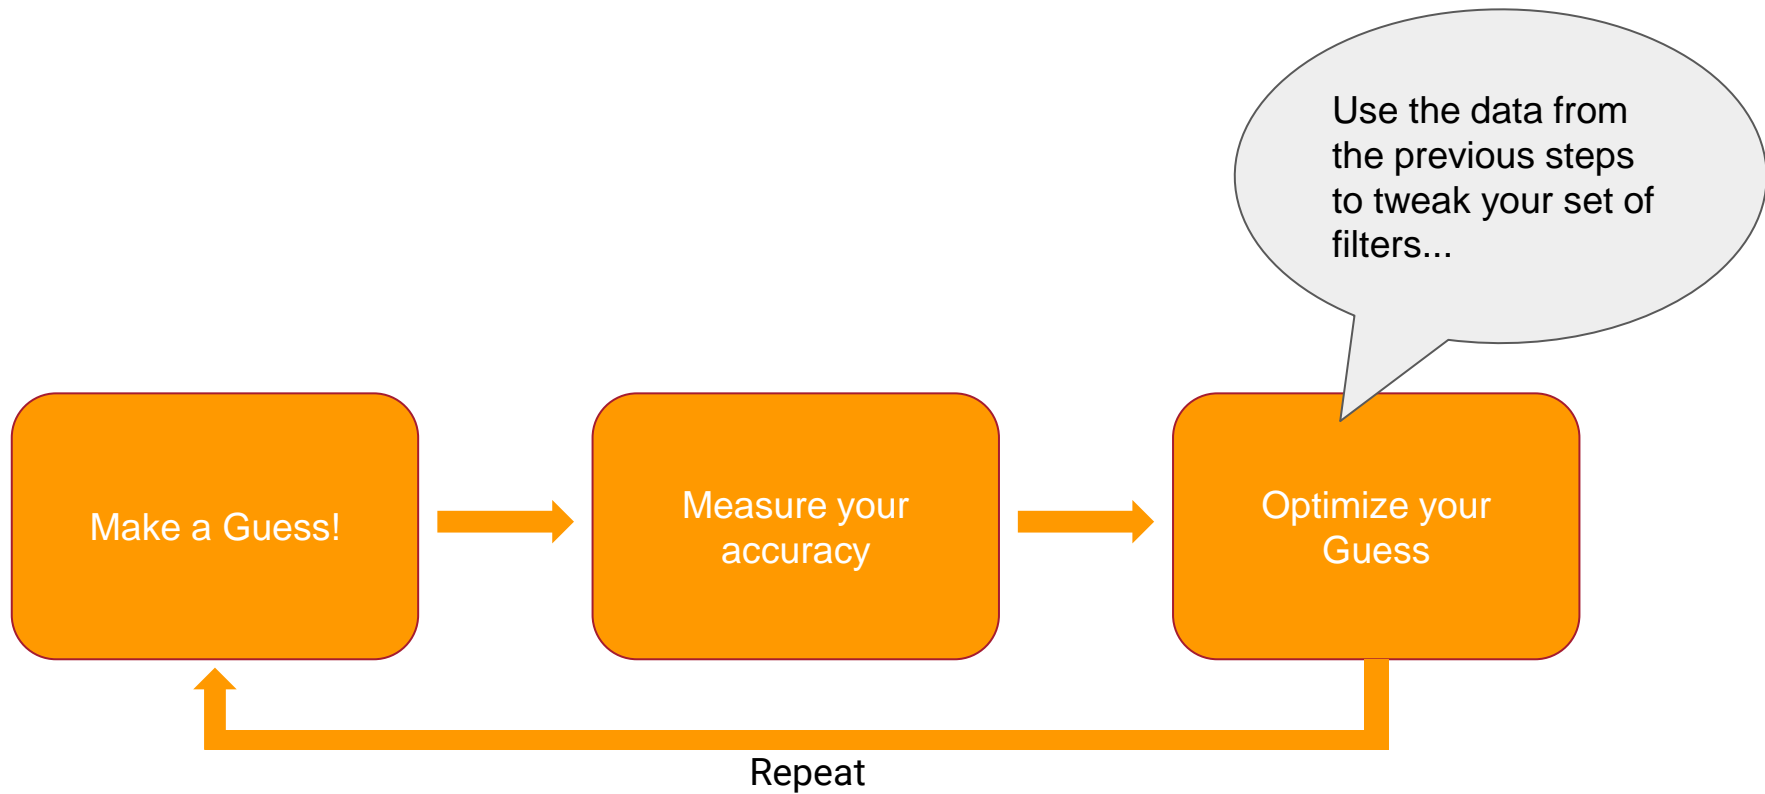


# The Machine Learning Paradigm



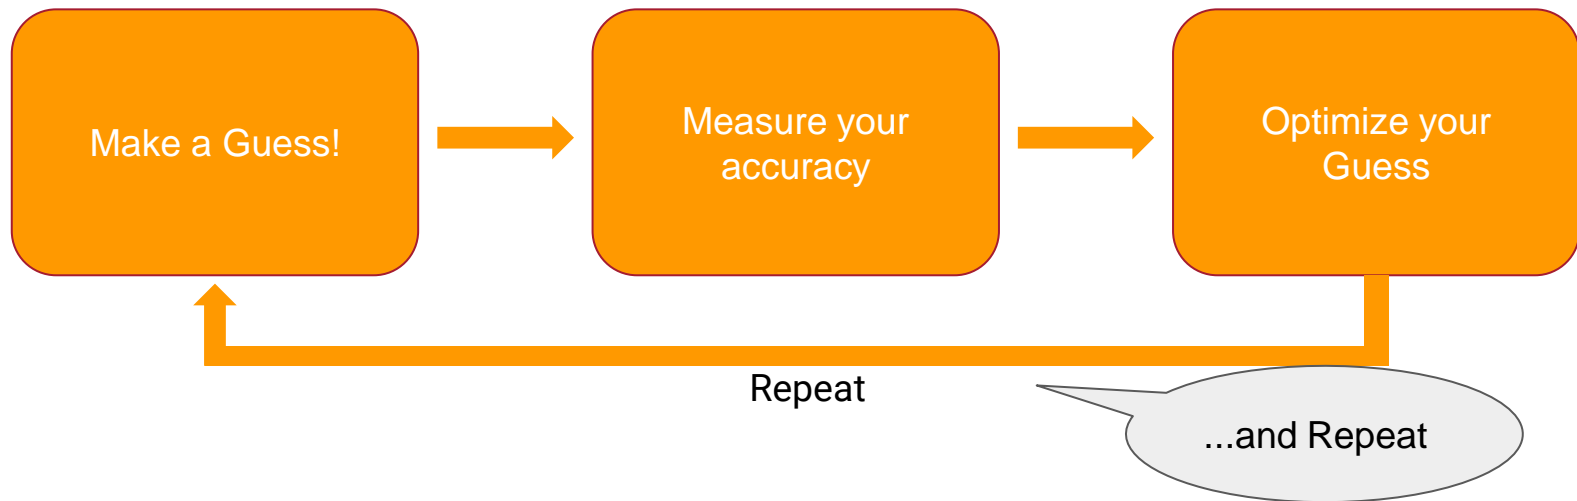




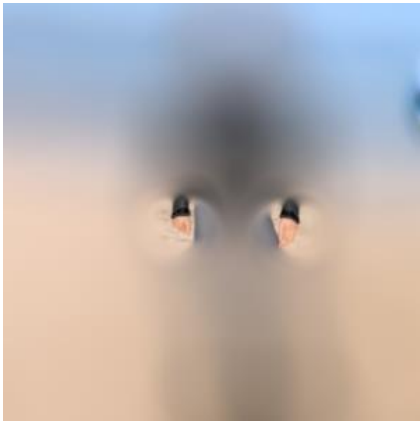
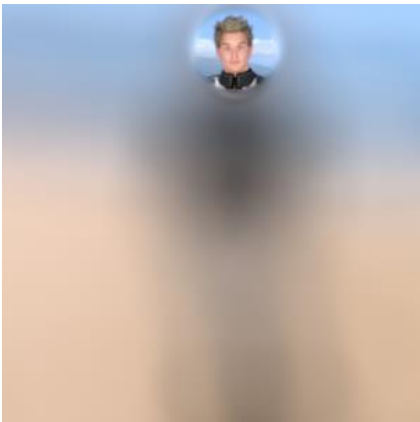

# The Machine Learning Paradigm






# The Machine Learning Paradigm




  +  +   = ?






+



+



=

HUMAN





```
model = tf.keras.models.Sequential([
    tf.keras.layers.Conv2D(16, (3, 3), activation='relu',
                           input_shape=(300, 300, 3)),
    tf.keras.layers.MaxPooling2D(2, 2),
    tf.keras.layers.Conv2D(32, (3, 3), activation='relu'),
    tf.keras.layers.MaxPooling2D(2, 2),
    tf.keras.layers.Conv2D(64, (3, 3), activation='relu'),
    tf.keras.layers.MaxPooling2D(2, 2),
    tf.keras.layers.Flatten(),
    tf.keras.layers.Dense(512, activation='relu'),
    tf.keras.layers.Dense(1, activation='sigmoid')
])
```

```
model = tf.keras.models.Sequential([
    tf.keras.layers.Conv2D(16, (3, 3), activation='relu',
                           input_shape=(300, 300, 3)),
    tf.keras.layers.MaxPooling2D(2, 2),
    tf.keras.layers.Conv2D(32, (3, 3), activation='relu'),
    tf.keras.layers.MaxPooling2D(2, 2),
    tf.keras.layers.Conv2D(64, (3, 3), activation='relu'),
    tf.keras.layers.MaxPooling2D(2, 2),
    tf.keras.layers.Flatten(),
    tf.keras.layers.Dense(512, activation='relu'),
    tf.keras.layers.Dense(1, activation='sigmoid')
])
```

Conv2D stands  
for 2D  
Convolution --  
another word for  
a filter

```
model = tf.keras.models.Sequential([
    tf.keras.layers.Conv2D(16, (3, 3), activation='relu',
                           input_shape=(300, 300, 3)),
    tf.keras.layers.MaxPooling2D(2, 2),
    tf.keras.layers.Conv2D(32, (3, 3), activation='relu'),
    tf.keras.layers.MaxPooling2D(2, 2),
    tf.keras.layers.Conv2D(64, (3, 3), activation='relu'),
    tf.keras.layers.MaxPooling2D(2, 2),
    tf.keras.layers.Flatten(),
    tf.keras.layers.Dense(512, activation='relu'),
    tf.keras.layers.Dense(1, activation='sigmoid')
])
```

MaxPooling is a way of compressing the image while enhancing features

```
model = tf.keras.models.Sequential([
    tf.keras.layers.Conv2D(16, (3, 3), activation='relu',
                           input_shape=(300, 300, 3)),
    tf.keras.layers.MaxPooling2D(2, 2),
    tf.keras.layers.Conv2D(32, (3, 3), activation='relu'),
    tf.keras.layers.MaxPooling2D(2, 2),
    tf.keras.layers.Conv2D(64, (3, 3), activation='relu'),
    tf.keras.layers.MaxPooling2D(2, 2),
    tf.keras.layers.Flatten(),
    tf.keras.layers.Dense(512, activation='relu'),
    tf.keras.layers.Dense(1, activation='sigmoid')
])
```


Dense is a  
neural network  
that matches the  
filters to the  
labels

```
model = tf.keras.models.Sequential([
    tf.keras.layers.Conv2D(16, (3, 3), activation='relu',
                           input_shape=(300, 300, 3)),
    tf.keras.layers.MaxPooling2D(2, 2),
    tf.keras.layers.Conv2D(32, (3, 3), activation='relu'),
    tf.keras.layers.MaxPooling2D(2, 2),
    tf.keras.layers.Conv2D(64, (3, 3), activation='relu'),
    tf.keras.layers.MaxPooling2D(2, 2),
    tf.keras.layers.Flatten(),
    tf.keras.layers.Dense(512, activation='relu'),
    tf.keras.layers.Dense(1, activation='sigmoid')
])
```


The final Dense  
represents the  
labels -  
Horse = 0,  
Human = 1



sigmoid




ReLU




```
optimizer = tf.keras.optimizers.RMSprop(lr=0.001)
model.compile(loss='binary_crossentropy',
              optimizer=optimizer,
              metrics=['accuracy'])
```

```
optimizer = tf.keras.optimizers.RMSprop(lr=0.001)
model.compile(loss='binary_crossentropy',
              optimizer=optimizer,
              metrics=['accuracy'])
```






```
optimizer = tf.keras.optimizers.RMSprop(lr=0.001)
model.compile(loss='binary_crossentropy',
              optimizer=optimizer,
              metrics=['accuracy'])
```

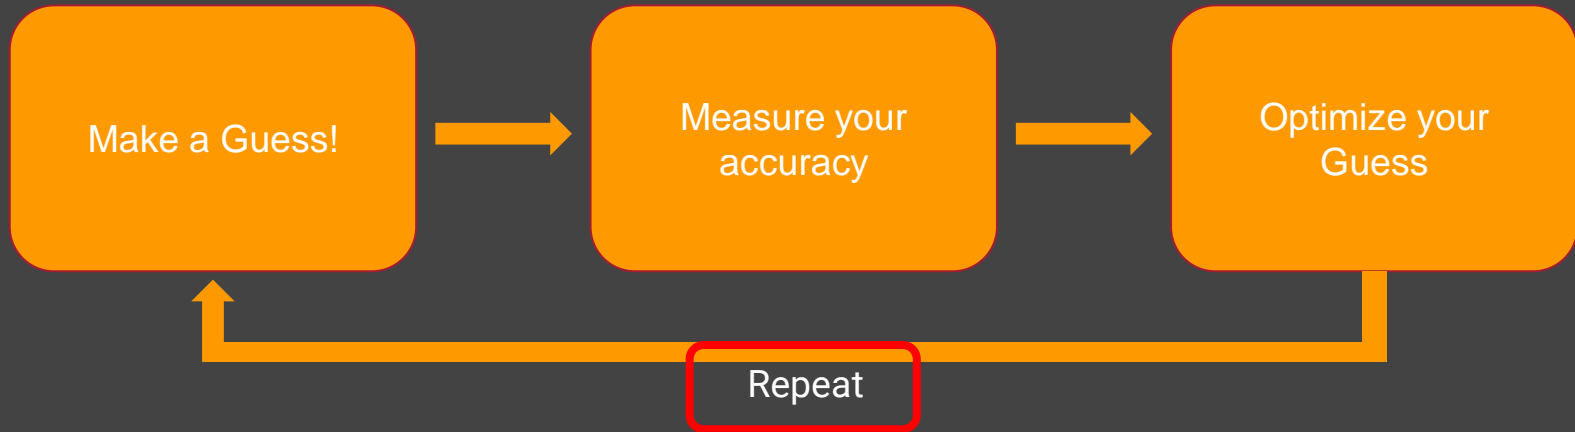


```
optimizer = tf.keras.optimizers.RMSprop(lr=0.001)
model.compile(loss='binary_crossentropy',
              optimizer=optimizer,
              metrics=['accuracy'])
```



```
model.fit(train_generator, epochs=12, ...)
```

```
model.fit(train_generator, epochs=12, ...)
```






# Demo

Training a computer to recognize Horses or Humans






Predicted: **wolf**  
True: **wolf**




Predicted: **husky**  
True: **husky**



Predicted: **wolf**  
True: **wolf**



Predicted: **wolf**  
True: **husky**



Predicted: **husky**  
True: **husky**



Predicted: **wolf**  
True: **wolf**

# Thanks !



Vinicius Fernandes Caridá

vfcarida@gmail.com



@vinicius caridá



@vfcarida



@vinicius caridá



@vfcarida



@vinicius caridá



@vfcarida